

Kartik Mahajan Associate

✉ kartikmahajan5688@gmail.com ☎ +91 9501145677 📍 India 🔗 LinkedIn

AI Engineer specializing in Generative AI and multi-agent systems, with experience building production conversational AI, RAG pipelines, and LLM-powered applications using AWS Strands, LangChain, Amazon Bedrock, and OpenAI. Proven track record of shipping AI solutions that improve accuracy and reduce manual effort.

Professional Experience

Cognizant Technology Solutions, Associate

February 2026 – Present | Gurugram, India

Project: Conversational Multi-Agent System for Healthcare Support

- Developed AI agents to interpret caller intent in real time, authenticate members, and route requests to specialist agents (Claims, ID Cards, Payment & Billing, and more).
- Built agent tooling via **Model Context Protocol (MCP)** and integrated Amazon Nova Sonic 2 for direct speech-to-speech caller interaction.
- Implemented **multi-agent orchestration** with AWS Strands on **Amazon Bedrock** and AWS Lambda, coordinating handoffs across the agent network.
- **Tech Stack:** Python · AWS Strands · Multi-Agent Systems · MCP · Amazon Bedrock · Amazon Nova Sonic 2 · AWS Lambda · Conversational AI · Intent Recognition

The Octavian Group 🔗, Application Developer

April 2025 – February 2026 | Remote, UK

- **Sober Meets:** Developed and deployed AI-driven predictive match analysis system for the Sober Meets app, Improved match accuracy by **35%** and user engagement by **20%** in three months by integrating **real-time feedback loops** using **LangChain** and **Pinecone**.
- **ZORQK:** Built a **conversational AI car-discovery chatbot** that parses natural-language queries into structured search filters using **LangChain**, improving search relevance for non-technical users.

Resolut Partners 🔗, Full Stack Developer

April 2023 – March 2025 | Mumbai, India

- Built an **AI-driven** time-tracking and billing platform (featuring an **admin dashboard, mobile, and desktop**), boosting **accuracy by 25%**.
- **Tech Stack - MongoDB, Express, React, Node JS**, React Native, Expo, Digital Ocean, AWS S3, Electron JS, Graph.js, Next.js, Material UI.

Certificates

- RAG and Agentic AI: Build Next-Gen AI Systems (Professional Certificate) 🔗
- Certified AI Practitioner — Amazon Web Services 🔗
- Generative AI Bedrock - Amazon Web Services 🔗

Education

Bachelor of Technology in Computer Science,

Lovely Professional University

July 2016 – June 2020 | Phagwara, India

Skills

Generative AI & LLMs

LLMs, Prompt Engineering, RAG, Multi-modal AI, Fine-tuning concepts, Hugging Face

Agentic AI & Frameworks

Multi-Agent Systems, AWS Strands, LangChain, Microsoft AutoGen, Model Context Protocol (MCP), Agent Orchestration, Tool/Function Calling

AI Infrastructure & Models

Amazon Bedrock, Amazon Nova Sonic 2, OpenAI, Vector Databases (Pinecone, ChromaDB), RAGAS (evaluation)

Programming & Backend

Python, FastAPI, JavaScript (ES6+), Node.js, Express.js, MongoDB

Cloud & Deployment

AWS (Bedrock, Lambda, S3), Azure, Docker, DigitalOcean

Web Development

React.js, React Native (Expo), Redux, Tailwind CSS, Material UI

Projects

AI-Powered Multi-Agent Blog Generator 🔗

- Built a full-stack multi-agent content pipeline (Researcher → Writer → Reviewer) using **FastAPI**, **LangChain**, and **OpenAI**, automating end-to-end blog creation and reducing manual content effort by 80%.

AI Mental Health Assistant 🔗

- Built a **multi-agent conversational AI** using **OpenAI** and **Microsoft AutoGen** to analyze user emotions and generate **therapy-based self-care suggestions**.

AI Document RAG Chatbot 🔗

- Built a **RAG chatbot** using **FastAPI**, **Gradio**, **LangChain**, and **Pinecone**, leveraging vector search to deliver context-aware responses and reduce hallucinations by 30%.